

Cascalog

Nathan Marz
Twitter
[@nathanmarz](https://twitter.com/nathanmarz)

Let's do some analysis



Tweets during the Tunisian revolution

What is Cascalog?

Abstraction



Cascalog

Variables and logic

Cascading

Tuples, data workflows

Hadoop

Key/value pairs,
simple aggregation

What is Hadoop MapReduce?

- High latency batch processing
- Massive scale (petabytes)
- Fault-tolerant

Why Cascalog?

Abstraction

Composition

Cascalog basics

```
(def age
  [["alice" 28]
   ["bob" 33]
   ["chris" 40]
   ["david" 25]
   ["emily" 25]
   ["george" 31]
   ["gary" 28]
   ["kumar" 27]
   ["luanne" 36]
  ])
```

The “age” dataset

Cascalog basics

```
(?<- (stdout) [?person] (age ?person ?age) (< ?age 30))
```


Where to
emit results

“Predicates”: constrain
the output variables

Output variables

Define and
execute a query

Predicates



Predicates

```
( * 2 ?x ::> ?z )
```

Fields can be constants or variables

Variables are prefixed with ? or !

Predicates

(+ 2 ?x :> 6)

(* 2 ?a :> ?z)
(* 3 ?b :> ?z)

(* ?x ?x :> ?x)

Predicates

- Functions
- Filters
- Aggregators
- Generators: finite sources of tuples

Example #1

```
(?<- (stdout) [?person] (age ?person ?age) (< ?age 30))
```



Generator



Filter

Example #2

```
(?<- (stdout) [?person]  
      (full-name ?person ?name) (extract-first-name ?name :> "Leon"))
```

Generator



Function



Example #3

```
(?<- (stdout) [?age]  
      (age ?person ?age) (c/count ?count) (> ?count 5))
```

Generator

Aggregator

Filter

Join example

```
(def follows  
  [["alice" "david"]  
   ["alice" "bob"]  
   ["alice" "emily"]  
   ["bob" "david"]  
   ["bob" "george"]])
```

```
(def gender  
  [["alice" "f"]  
   ["bob" "m"]  
   ["chris" "m"]  
   ["david" "m"]  
   ["emily" "f"]])
```

```
(?<- (stdout) [?person]  
      (follows "emily" ?person) (gender ?person "m"))
```

Triggers a join



Join example

```
(def follows  
  [["alice" "david"]  
   ["alice" "bob"]  
   ["alice" "emily"]  
   ["bob" "david"]  
   ["bob" "george"]])
```

```
(def gender  
  [["alice" "f"]  
   ["bob" "m"]  
   ["chris" "m"]  
   ["david" "m"]  
   ["emily" "f"]])
```

```
(?<- (stdout) [?person]  
      (follows "emily" ?person) (gender ?person "m"))
```

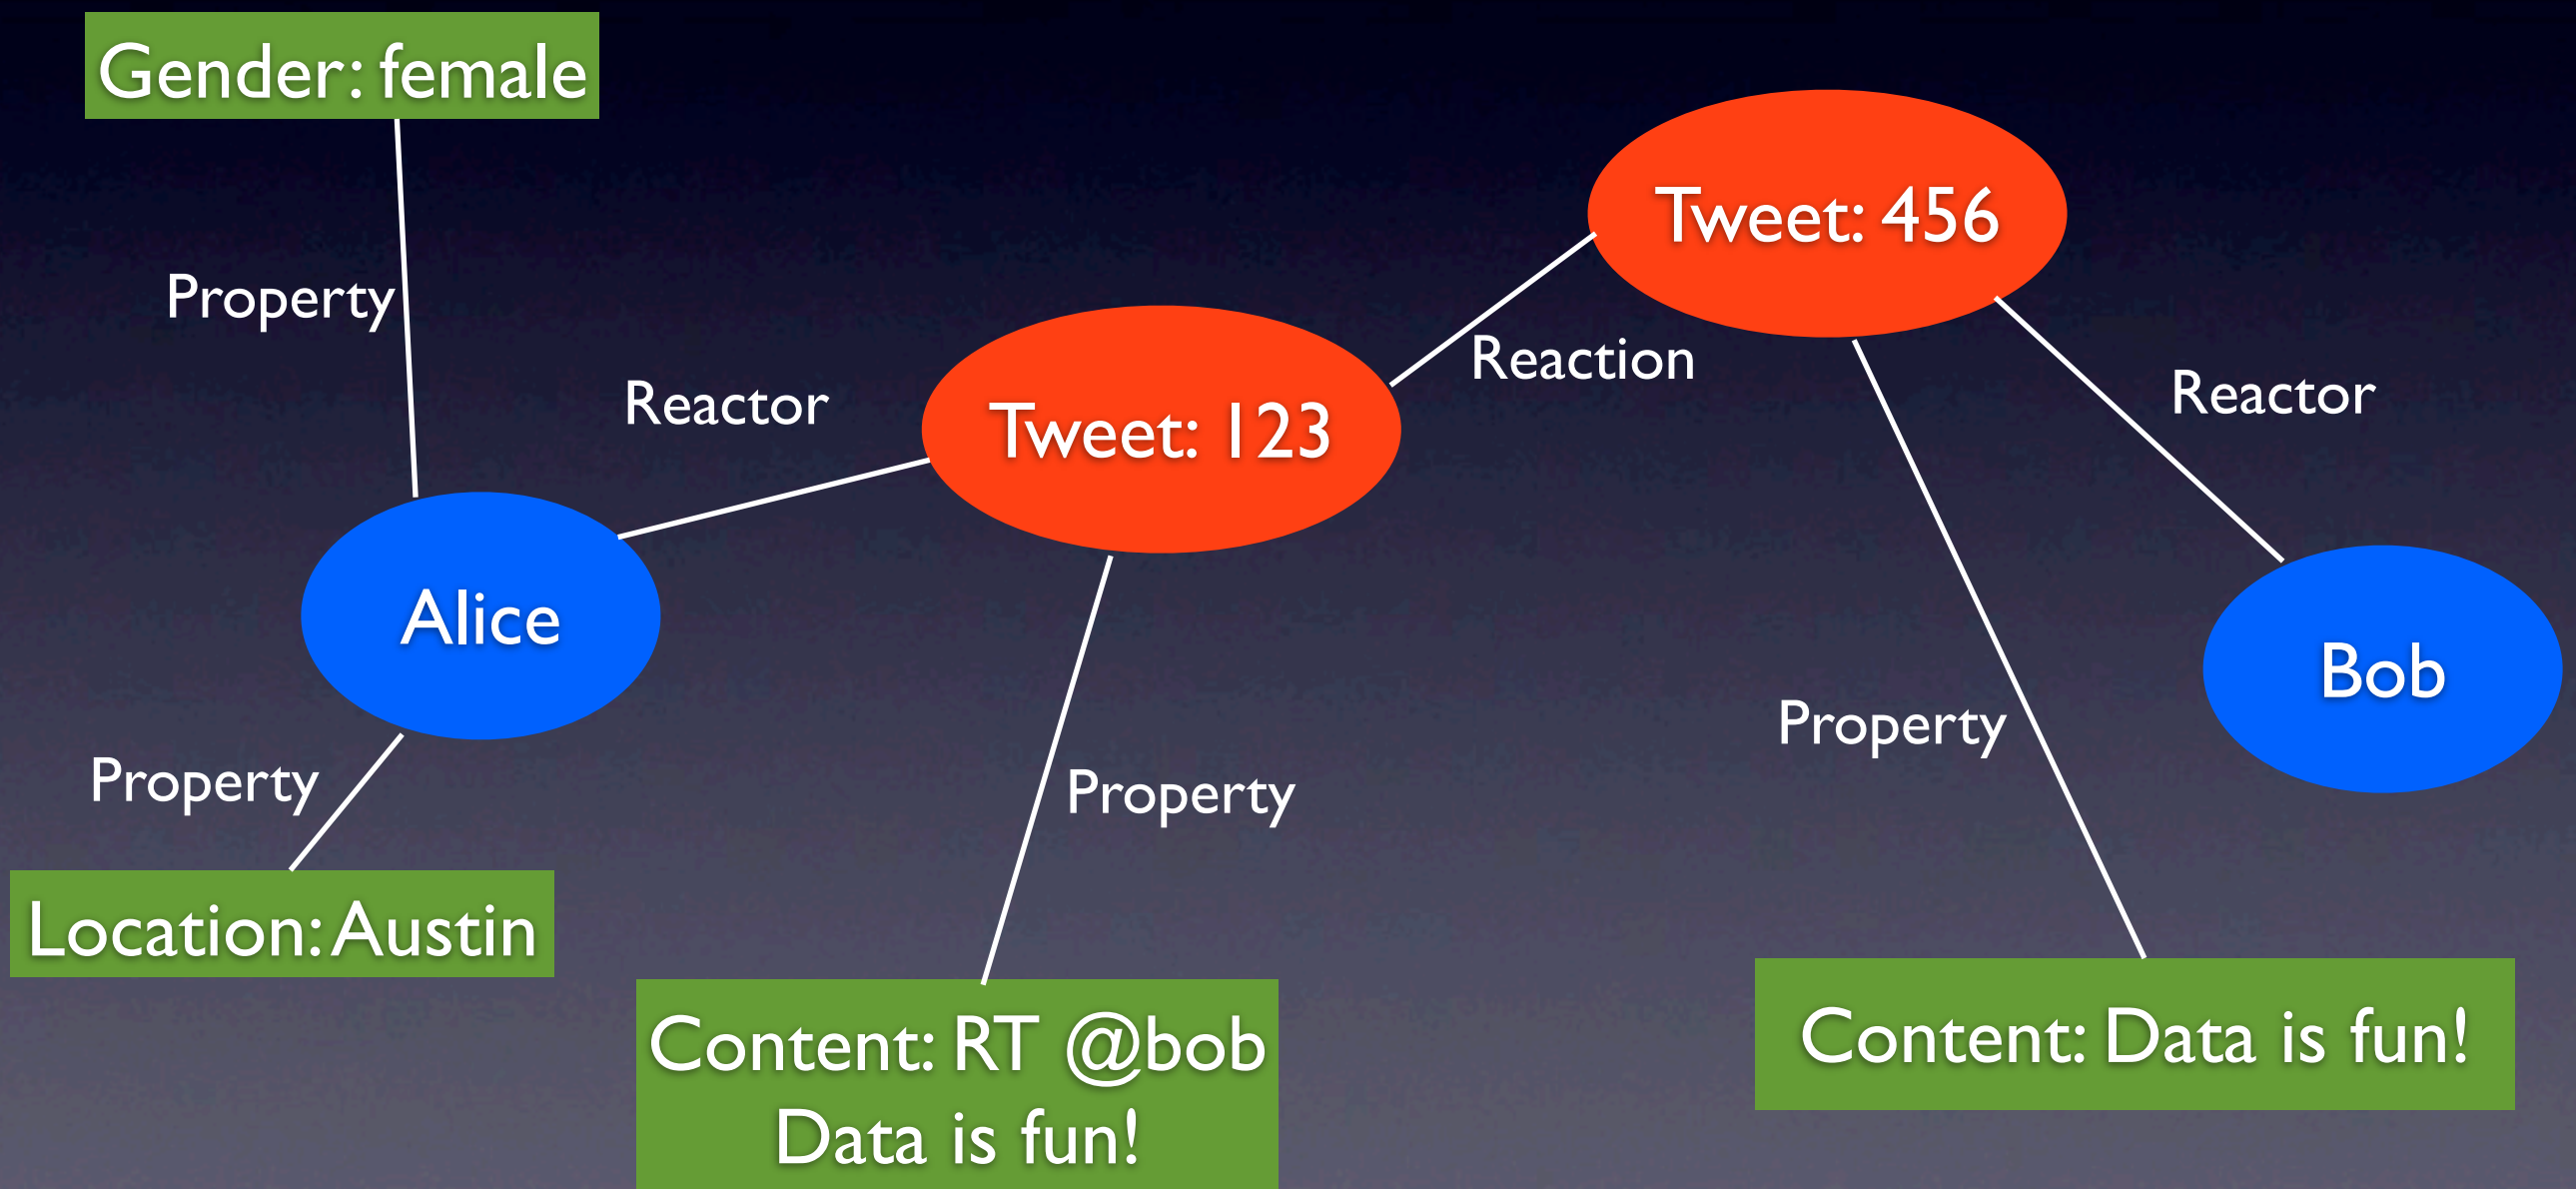
Joins are implicit

Demo code

The screenshot shows the GitHub interface for the repository 'nathanmarz / cascalog-conj'. At the top, the GitHub logo and 'SOCIAL CODING' tagline are on the left. The user 'nathanmarz' is logged in, with links to 'Dashboard', 'Inbox 0', 'Account Settings', and 'Log Out'. Below this is a navigation bar with 'Explore GitHub', 'Gist', 'Blog', 'Help', and a search bar. The repository header shows the name 'nathanmarz / cascalog-conj' and action buttons: 'Admin', 'Unwatch', 'Fork', 'Pull Request', and statistics for watchers (1) and forks (1). A tabbed interface below the header includes 'Code' (selected), 'Network', 'Pull Requests 0', 'Issues 0', 'Wiki 0', and 'Stats & Graphs'. The 'Code' tab displays a description: 'Code from my presentation of Cascalog at Clojure/conj 2011 — Read more' and a link to 'http://cascalog.org'. Below this are cloning options: 'Clone in Mac', 'ZIP', 'SSH', 'HTTP', 'Git Read-Only', and the repository URL 'git@github.com:nathanmarz/cascalog-conj.git'. A 'Read+Write access' badge is also present. Further down, a 'Files' tab is active, showing a list of files. The 'Commits' tab shows the latest commit to the 'master' branch, 'more examples', authored by 'nathanmarz' just now, with commit hash '6a36059fdd'. At the bottom, a table lists the files in the repository.

name	age	message	history
src/	just now	more examples [nathanmarz]	

Data model



Composability

```
(def avg  
  (<- [<val :> !avg]  
      (c/count !count)  
      (c/sum !val :> !sum)  
      (div !sum !count :> !avg)))
```

“Predicate macro”


Composability

```
(?<- (stdout) [?avg-age]  
      (age _ ?age)  
      (avg ?age :> ?avg-age))
```

expands to

```
(?<- (stdout) [?avg-age]  
      (age _ ?age)  
      (c/count !count)  
      (c/sum ?age :> !sum)  
      (div !sum !count :> ?avg-age))
```


Attaching chains

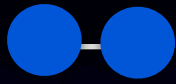





Chains

Chain

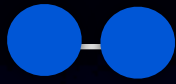




Chains

Chain				
Reverse binary (length - 1)	0	1	1	

Chain of length 8

Chains

Chain				
Reverse binary (length - 1)	0	1		

Chain of length 3

Chains

Chain	2	3	5	9	17	33
Reverse binary (length - 1)	1	0	1	0	1	

Chain of length 22

Questions

<http://github.com/nathanmarz/casalog>